



Probabilistic Programming With PyMC3

John Kimball

9/26/19




Overview

- Why do we need Probabilistic Programming (PP)?
- What is PP?
- Four simple examples of PP using PyMC3
 - A bit of theory as needed
- How PP is being used today
- How to get started

The goal is a high-level overview, not a detailed tutorial



Why Do We Need Probabilistic Programming (PP)?

- An explosion of data
 - Web commerce, Internet of Things, human genome, etc.
 - Often the data is uncertain
 - Enormous value is theoretically available from the data
 - Very difficult to interpret the data (draw inferences)
 - Advanced and specialized mathematical knowledge is required
 - Inference methods are unique to each problem
 - It is difficult to experiment and make changes to the inference systems
 - See video:
 - <https://www.oreilly.com/ideas/probabilistic-programming>
- 




Probabilistic Programming to the Rescue

- Two major components of a PP system:
 - **The modeling system is a high level programming language (e.g. Python)**
 - All the flexibility of the host language for expressing data structures
 - **Probability distributions as high level constructs**
 - **An automatic general purpose inference engine**
 - Independent from the model
 - Enables “running the model backward” to infer causes from data
- The PP system provides:
 - High-level probability constructs => shorter code, reduced development time
 - Host language flexibility => more sophisticated models
 - Automatic inference => greatly reduced statistical expertise, rapid experimentation
- Goal:
 - Domain experts can learn from their own data
 - Statistics/Machine Learning experts can increase their power



Component 1: Probabilistic Modeling in PyMC3

- Model setup
 - Probability distributions
 - Continuous: Normal Distribution
 - Discrete: Bernoulli Distribution
 - Sampling data from the model
 - Inspecting model output using Arviz
 - See Jupyter Notebook:
 - `Part1-ProbabilisticModeling.ipynb`
- 

Running Models in Reverse

- Bayes' Theorem:

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

$P(A|B)$ means the probability of A being true given that B is true

- Bayes' Theorem gives us a way of using data to update our models

Example: Ferrari Detector

- ▶ We are designing a sensor to detect whether a car is a Ferrari. Our computer processor can't run computer vision software, but can detect color. If our sensor detects red, how likely is it that the car is a Ferrari?

- ▶ We know that

- ▶ 50% of Ferraris are red
- ▶ 5% of all cars are red
- ▶ .02% of all cars are Ferraris

Let:

R be True if the car is red

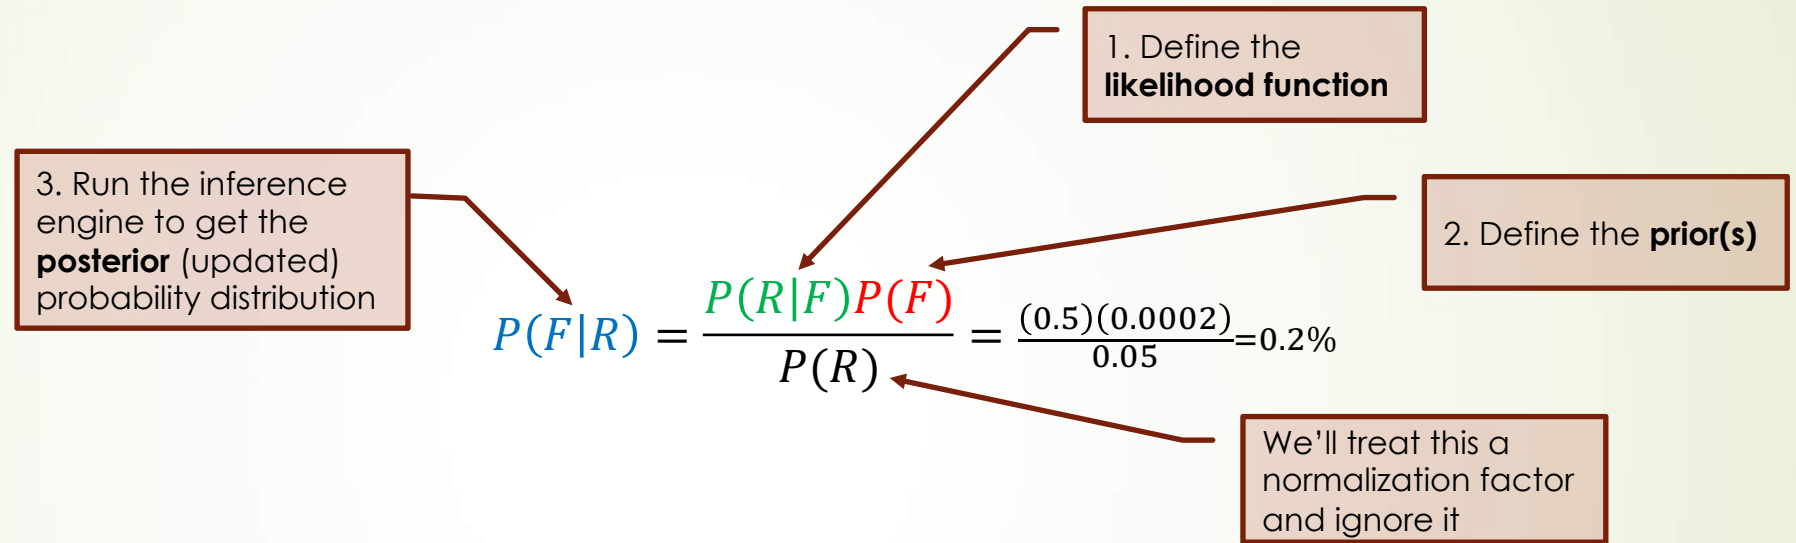
F be True if the car is a Ferrari

We want to know the probability a passing car is a Ferrari if all we know is that it is red.

$$P(F|R) = \frac{P(R|F)P(F)}{P(R)} = \frac{(0.5)(0.0002)}{0.05} = 0.2\%$$

- ▶ **Prior** to any evidence, there is a 0.02% probability of a car being Ferrari
- ▶ The **likelihood** of a data (Red) given the hypothesis (Ferrari)
- ▶ After (**Posterior** to) the evidence that the car is red, there is a 0.2% probability that the car is a Ferrari

Terminology and Procedure¹



- **Prior** to any evidence, there is a 0.02% probability of a car being Ferrari
- The **Likelihood** of data (Red) given the hypothesis (Ferrari)
- After (**Posterior** to) the evidence that the car is red, there is a 0.2% probability that the car is a Ferrari
- Note that in most models:
 - The the priors, likelihoods and posterior will be in terms of *probability distributions*
 - The priors will be in terms of variable parameters, and the results will include probability distributions for the parameters, hence updating the model



Coin Toss

- The "Hello, World!" of probabilistic programming
- We'll:
 - Generate some data with known bias (outside of PyMC3)
 - Create a PyMC3 model of the coin toss process
 - Pass the data to our model
 - Attempt to determine the bias in the coin
- See Jupyter Notebook:
 - `Part2-FairCoin.ipynb`



A/B Testing

- Given:
 - Two versions of a website (A and B) differing in a single feature
 - We expect B to be an improvement on the original A
- Then:
 - Randomly present A and B to different groups
 - Observe which users purchase the product
 - Use PP to infer the effectiveness of the change from A to B
- See Jupyter Notebook:
 - Part3-ABTesting.ipynb



Challenger Disaster

- Given:
 - O-ring damage and temperature for 23 launches
- Then:
 - Model the probability of O-Ring damage vs. temperature
 - Infer model parameters from data
 - Display the expected probability of damage
 - Include 95% confidence intervals
 - Evaluate probability of damage at Challenger launch (31F)
- See Jupyter Notebook:
 - `Part4-Challenger.ipynb`

Inference Algorithms

- Non-Markovian
 - Brute force: Grid method
 - Define a grid on which to evaluate each parameters
 - Multiply the prior and the likelihood at each point to define the posterior at that point
 - Very inefficient; spends much time in areas that don't matter
 - ADVI
 - Approximates posterior as a simpler function
 - Becomes an optimization problem
 - Uses automatic differentiation gradients for efficient optimization
- Markov Chain Monte Carlo
 - Evaluates the prior and likelihood function pointwise (like grid method)
 - Much more efficient than grid method, since spends time in most important areas

See Martin, Osvaldo, *Bayesian Analysis with Python: Introduction to Statistical Modeling and Probabilistic Programming Using PyMC3 and ArviZ*, Chapter 8

Other Areas Probabilistic Programming is Being Used

- PyMC3
 - According to Thomas Wiecki, A | B testing is the most common use case
 - Quantitative Finance (e.g. Quantopian to evaluate trading algorithms)
 - Earthquake Analysis (e.g. BEAT <https://github.com/hvasbath/beat>)
 - Supply Chain Optimization (https://twiecki.io/blog/2019/01/14/supply_chain/)
- Probabilistic Programming in General
 - Astrophysics (used to infer black hole parameters in support of 2017 Nobel Prize in Physics)
 - Search and Rescue
 - https://en.wikipedia.org/wiki/Search_and_Rescue_Optimal_Planning_System
 - Drug Discovery
 - "...allows for predicting complex phenotypes by casting light on the causal molecular underpinnings of disease. This method combines the use of **ensemble deep learning with probabilistic programming that applies Bayesian models to test the causal dependencies** of millions of possible interactions between proteins to elucidate the core biological network connecting the observed phenotype with the mechanism of disease."
 - <https://www.wuxinextcode.com/genomic-insights/a-i-breakthroughs-in-drug-rd-probabilistic-programming-biological-context-and-quantum-a-i/>

There are countless other applications of probabilistic programming

More Python Libraries for Probabilistic Programming

- PyStan – Python interface to the popular Stan PP platform
- Edward – “a Python library for probabilistic modeling, inference, and criticism. ... Edward fuses three fields: Bayesian statistics and machine learning, deep learning, and probabilistic programming.
- Pyro (Uber AI Labs) – “Pyro is a universal probabilistic programming language (PPL) written in Python ... enables flexible and expressive deep probabilistic modeling, unifying the best of modern deep learning and Bayesian modeling.
- MIT Probabilistic Computing Project - <http://probcomp.csail.mit.edu/software/>
- PyCBC – specialized for LIGO, which won 2017 Nobel Prize for Physics
 - “PyCBC was used in the first direct detection of gravitational waves (GW150914) by LIGO and is used in the ongoing analysis of LIGO and Virgo data.”
<http://pycbc.org/pycbc/latest/html/>
 - MCMC Bayesian inference like PyMC3
<http://pycbc.org/pycbc/latest/html/inference.html>

References

- What is probabilistic programming?
 - <https://www.oreilly.com/ideas/probabilistic-programming>
- PyMC3 Online Documentation
 - <https://pymc3.readthedocs.io/en/latest/index.html>
- Martin, Osvaldo, *Bayesian Analysis with Python: Introduction to Statistical Modeling and Probabilistic Programming Using PyMC3 and ArviZ*, 2nd edition, Packt Publishing, 2018.
- Davidson-Pilon, Cameron, *Bayesian Methods for Hackers: Probabilistic Programming and Bayesian Inference*, Addison Wesley, 2016.
 - Also available online:
 - <https://camdavidsonpilon.github.io/Probabilistic-Programming-and-Bayesian-Methods-for-Hackers/>
 - <https://github.com/CamDavidsonPilon/Probabilistic-Programming-and-Bayesian-Methods-for-Hackers>
- The Python Podcast.__init__ (April 29, 2019 with guest Thomas Wiecki)
 - <https://www.podcastinit.com/feed/mp3/>
- Pfeffer, Avi, *Practical Probabilistic Programming*, Manning Publications, 2016.
 - Uses Figaro, a Scala-based PP system.

Additional Learning Resources

- Kruschke, John K., *Doing Bayesian Data Analysis: A Tutorial with R and BUGS*, Academic Press, 2011.
 - PyMC3 port of examples:
 - https://github.com/aloctavodia/Doing_bayesian_data_analysis
- Thomas Wiecki (PyMC3 developer) blog:
 - <https://twiecki.io>
- McGrayne, Sharon Bertsch, *The Theory That Would Not Die: How Bayes' Rule Cracked the Enigma Code, Hunted Down Russian Submarines, and Emerged Triumphant from Two Centuries of Controversy*, Yale University Press, 2011.
- VanderPlas, Jake, *Frequentism and Bayesianism: A Practical Introduction*
 - <http://jakevdp.github.io/blog/2014/03/11/frequentism-and-bayesianism-a-practical-intro>

How to install (the easy way)

- 1. Make sure Docker is installed
- 2. From a terminal window, get a docker image from dockerhub:
 - `docker pull scalafan/pymc3-arviz:version_1`
- 3. From a terminal window, run scalafan/pymc3_arviz:
 - `docker run -p 8888:8888 -v /Users/yourhome/PyMC3Models:/home/jovyan scalafan/pymc3-arviz:version_1`
- 4. Copy and paste the provided URL into your browser address bar
 - Any notebooks you create will be saved in the PyMC3Models directory
 - The host directory (PyMC3Models) must be shared in Docker, and user jovyan must have read and write privileges
- 5. Enter Ctrl-C in the terminal window to shut down the notebook server
- 6. You may wish to clean up the stopped Docker container:
 - `docker ps -a` # this will give a container name
 - `docker rm containername`

<https://hub.docker.com/r/scalafan/pymc3-arviz>



Additional Information



The Python Data Science Stack

- A great overview by Jake VanderPlas (though a bit dated)
 - <https://speakerdeck.com/jakevdp/pythons-data-science-stack-jsm-2016>
- Some other libraries used by PyMC3 and this presentation:
 - Theano: a Python library that lets you to define, optimize, and evaluate mathematical expressions
 - Automatically generated C code
 - Automatic differentiation of mathematical expressions
 - PyMC4 will use TensorFlow Probability instead of Theano
 - Theano is being maintained by the PyMC team
 - Arviz: a Python package for exploratory analysis of Bayesian models.
 - Much easier plotting than Matplotlib

Examples of ML Problems in Probabilistic Programming

- ▶ Bayesian Linear Regression
- ▶ Logistic Regression
- ▶ Naïve Bayes
- ▶ K-Means Clustering
- ▶ Latent Dirichlet Allocation (LDA)
- ▶ Correlated Topic Models (CTM)
- ▶ Autoregressive Integrated Moving Average (ARIMA)
- ▶ Hidden Markov Models
- ▶ Matrix Factorization
- ▶ Sparsity and Sparse Bayes
- ▶ Conditional Random Fields

PP makes many machine learning tasks available to a broad audience